


Introduction to Recurrent Neural Networks – Part 1

Devashish Khatwani
June 2018



A Motivating example for sequence modelling: Predict the next word

"This morning I took the dog for a walk."



"This morning I took the dog for a walk."

given these words

predict what comes next?



"This morning I took the dog for a walk."

given these 2 words, predict the next word



[1 0 0 0 0 0 1 0 0 0]

for

a

One hot feature vector indicates what each word is

prediction

But...

“In **France**, I had a great time and I learnt some of the _____
language.”

We need information from the far past and future to accurately guess the correct word.

Try using the whole sentence as a window!

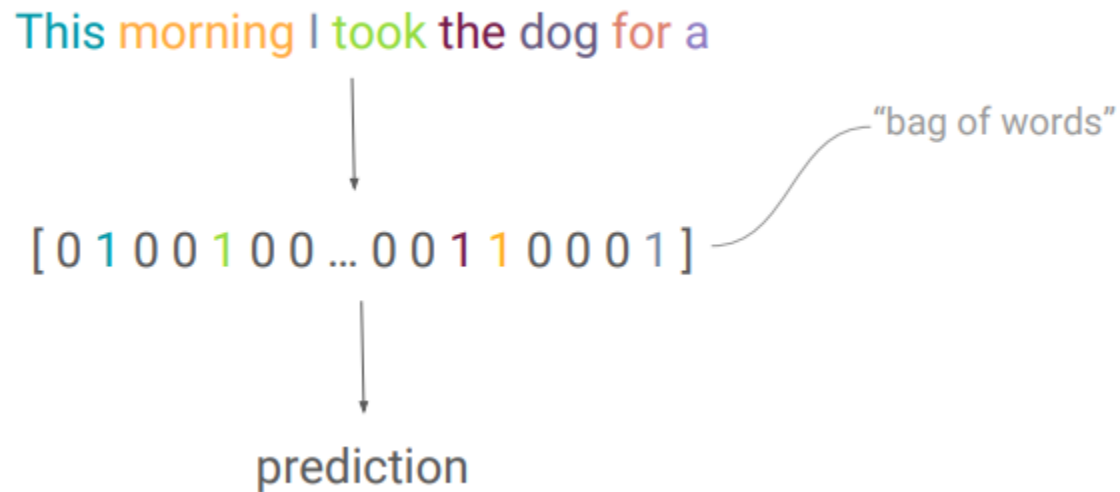
This morning I took the dog for a

[0 1 0 0 1 0 0 ... 0 0 1 1 0 0 0 1]

"bag of words"

prediction

Try using the whole sentence as a window!



But... Counts don't preserve order

"The food was good, not bad at all."
VS

"The food was bad, not good at all."

Use really big windows!

"This morning I took the dog for a walk."
given these 7 words, predict the next word

[1 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 ...]

morning I took the dog ...



↓
prediction



But... Curse of Dimensionality

Difference with CNNs


- Convolution in space (CNN) VS convolution in time (RNN)
- CNN: models relationships in space. Filter slides along x and y dimensions
- RNN: models relationships in time. “Filter” slides along time dimension



What is a Sequence?


1

Audio



2

ECG Reading



3

Sentence

Lucy is going to the park.

4

Bank Data?

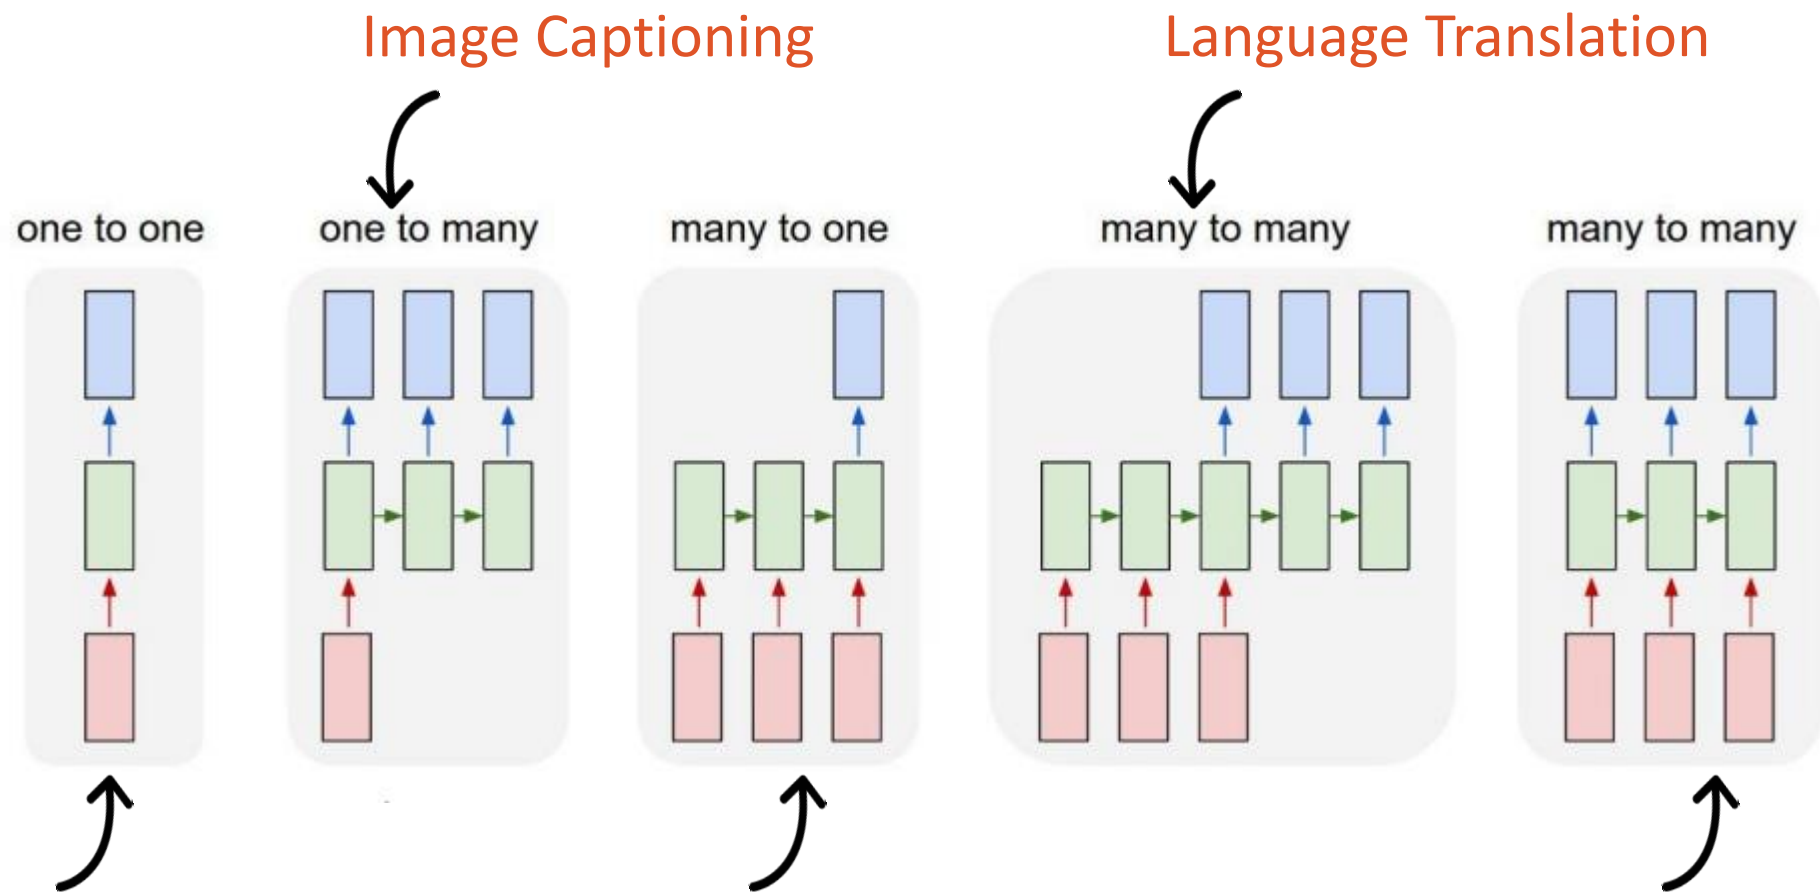
To solve our problem we need...

1. To deal with **variable-length** sequences
2. To maintain **sequence order**
3. To keep track of **long-term dependencies**
4. To **share parameters** across the sequence



Try out Recurrent Neural Networks

Types of RNNs




Not really a RNN!

Sentiment Classification

Video classification at frame level

RNNs remember their previous state




x_0 : vector representing first word
 s_0 : cell state at $t = 0$ (some initialization)
 s_1 : cell state at $t = 1$

$$s_1 = \tanh(Wx_0 + Us_0)$$

W, U : weight matrices

RNNs remember their previous state



x_1 : vector representing second word


s_1 : cell state at $t = 1$

s_2 : cell state at $t = 2$

$$s_2 = \tanh(Wx_1 + Us_1)$$


W, U : weight matrices

Unfolding the RNN across time



Notice that W and U are shared parameters across time


Training a RNN: What is our Loss?



Training a RNN: What are our Gradients?


we sum gradients across time for each parameter P :

$$\frac{\partial J}{\partial P} = \sum_t \frac{\partial J_t}{\partial P}$$



$$\frac{\partial J}{\partial W} = \sum_t \frac{\partial J_t}{\partial W}$$

Training a RNN: Try it out with W




$$\frac{\partial J}{\partial W} = \sum_t \frac{\partial J_t}{\partial W}$$

so let's take a single timestep t :

$$\frac{\partial J_2}{\partial W} = \frac{\partial J_2}{\partial y_2}$$

Training a RNN: Try it out with W




$$\frac{\partial J}{\partial W} = \sum_t \frac{\partial J_t}{\partial W}$$

so let's take a single timestep t :

$$\frac{\partial J_2}{\partial W} = \frac{\partial J_2}{\partial y_2} \frac{\partial y_2}{\partial s_2}$$

Training a RNN: Try it out with W




$$\frac{\partial J}{\partial W} = \sum_t \frac{\partial J_t}{\partial W}$$

so let's take a single timestep t :

$$\frac{\partial J_2}{\partial W} = \frac{\partial J_2}{\partial y_2} \frac{\partial y_2}{\partial s_2} \frac{\partial s_2}{\partial W}$$

Training a RNN: Try it out with W



$$\frac{\partial J}{\partial W} = \sum_t \frac{\partial J_t}{\partial W}$$


so let's take a single timestep t :

$$\frac{\partial J_2}{\partial W} = \frac{\partial J_2}{\partial y_2} \frac{\partial y_2}{\partial s_2} \frac{\partial s_2}{\partial W}$$

but wait...

$$s_2 = \tanh(U s_1 + W x_2)$$

Training a RNN: Try it out with W



$$\frac{\partial J}{\partial W} = \sum_t \frac{\partial J_t}{\partial W}$$

so let's take a single timestep t :

$$\frac{\partial J_2}{\partial W} = \frac{\partial J_2}{\partial y_2} \frac{\partial y_2}{\partial s_2} \frac{\partial s_2}{\partial W}$$


but wait...

$$s_2 = \tanh(U s_1 + W x_2)$$

s_1 also depends on W so we can't just treat $\frac{\partial s_2}{\partial W}$ as a constant!

Training a RNN: Try it out with W

how does s_2 depend on W ?



$$\begin{aligned} & \frac{\partial s_2}{\partial W} \\ & + \frac{\partial s_2}{\partial s_1} \frac{\partial s_1}{\partial W} \\ & + \frac{\partial s_2}{\partial s_0} \frac{\partial s_0}{\partial W} \end{aligned}$$

Training a RNN: Backpropagation Through Time

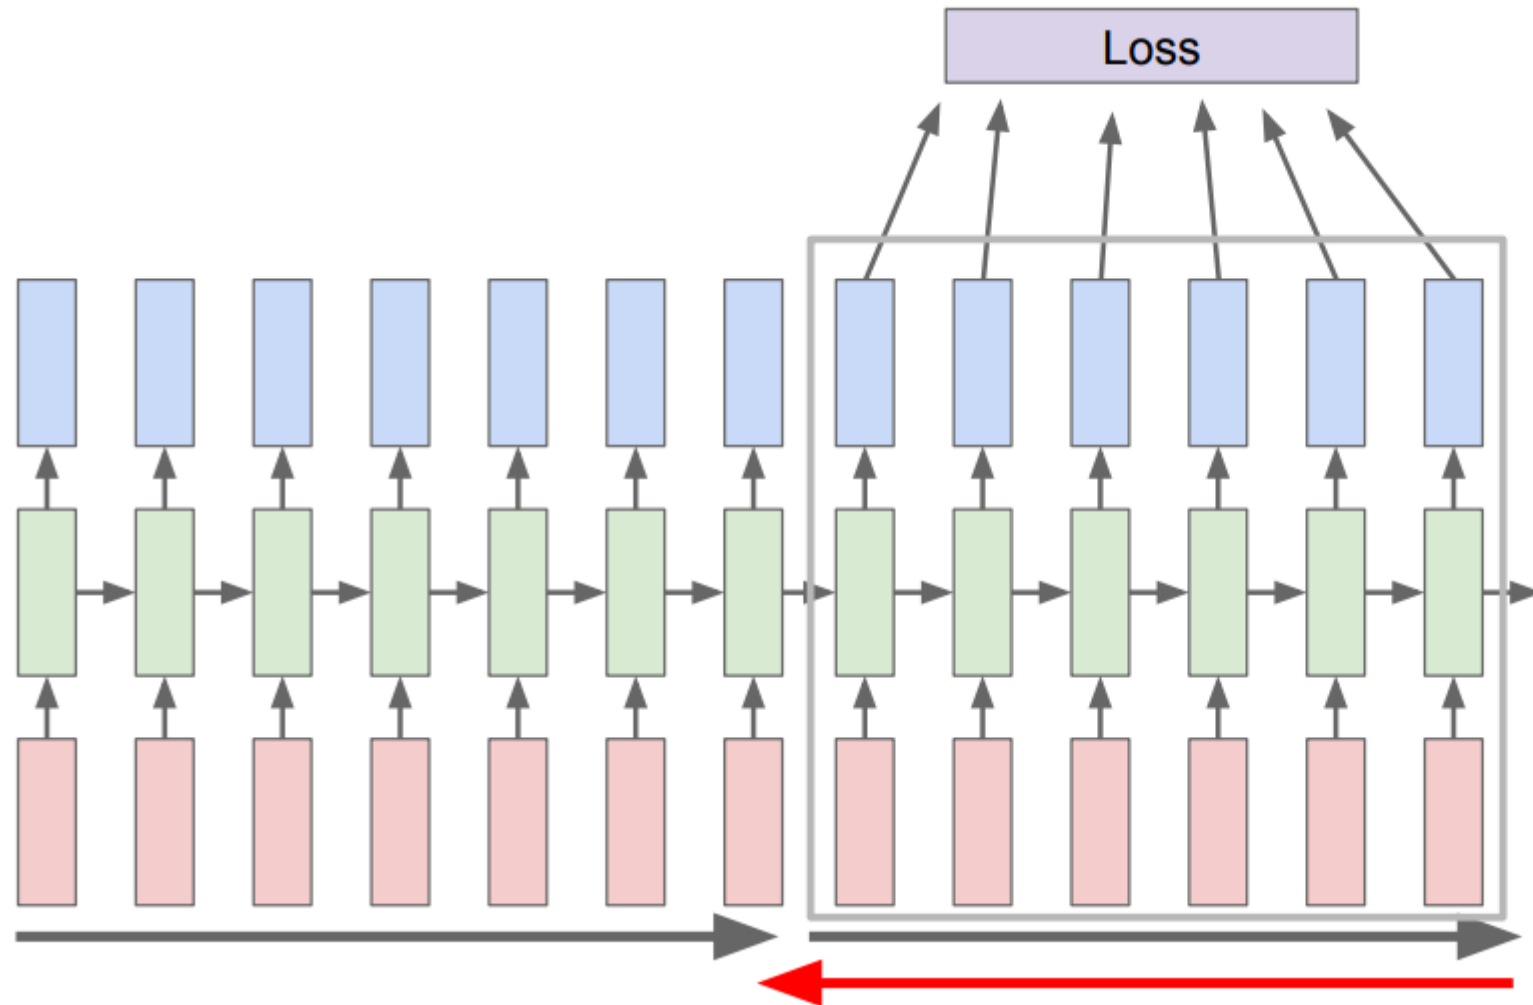
$$\frac{\partial J_2}{\partial W} = \sum_{k=0}^2 \frac{\partial J_2}{\partial y_2} \frac{\partial y_2}{\partial s_2} \underbrace{\frac{\partial s_2}{\partial s_k} \frac{\partial s_k}{\partial W}}_{\text{Contributions of } W \text{ in previous timesteps to the error at timestep } t}$$

Contributions of W in previous timesteps to the error at timestep t

$$\frac{\partial J_t}{\partial W} = \sum_{k=0}^t \frac{\partial J_t}{\partial y_t} \frac{\partial y_t}{\partial s_t} \underbrace{\frac{\partial s_t}{\partial s_k} \frac{\partial s_k}{\partial W}}_{\text{Contributions of } W \text{ in previous timesteps to the error at timestep } t}$$

Contributions of W in previous timesteps to the error at timestep t

Training a RNN: Truncated Backpropagation Through Time




Carry hidden states forward in time forever, but only backpropagate for some smaller number of steps

But RNNs are hard to train: Vanishing Gradient


$$\frac{\partial J_2}{\partial W} = \sum_{k=0}^2 \frac{\partial J_2}{\partial y_2} \frac{\partial y_2}{\partial s_2} \frac{\partial s_2}{\partial s_k} \frac{\partial s_k}{\partial W}$$

at $k = 0$:

$$\frac{\partial s_2}{\partial s_0} = \frac{\partial s_2}{\partial s_1} \frac{\partial s_1}{\partial s_0}$$



$$\frac{\partial J_n}{\partial W} = \sum_{k=0}^n \frac{\partial J_n}{\partial y_n} \frac{\partial y_n}{\partial s_n} \frac{\partial s_n}{\partial s_k} \frac{\partial s_k}{\partial W}$$



Two Hacky Solutions

what are each of these terms? →

$$\frac{\partial s_n}{\partial s_{n-1}} \frac{\partial s_{n-1}}{\partial s_{n-2}} \cdots \frac{\partial s_3}{\partial s_2} \frac{\partial s_2}{\partial s_1} \frac{\partial s_1}{\partial s_0}$$

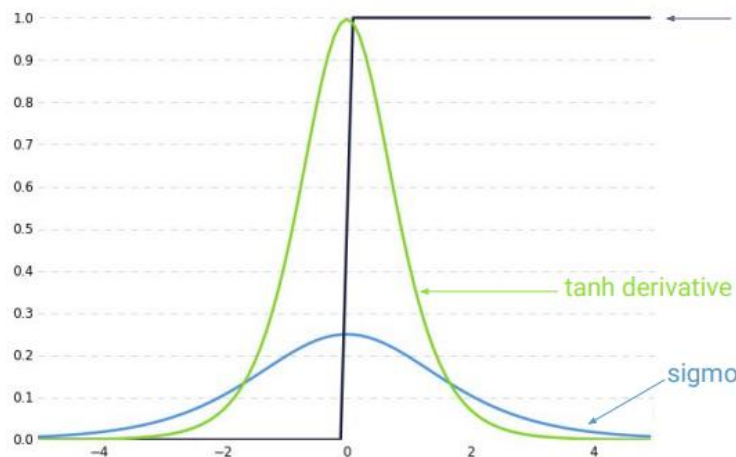
$$\frac{\partial s_n}{\partial s_{n-1}} = W^T \text{diag}[f'(W s_{j-1} + U x_j)]$$

W = sampled from standard normal distribution = mostly < 1

f = tanh or sigmoid so $f' < 1$



1



ReLU derivative

prevents f' from shrinking the gradients

tanh derivative

sigmoid derivative

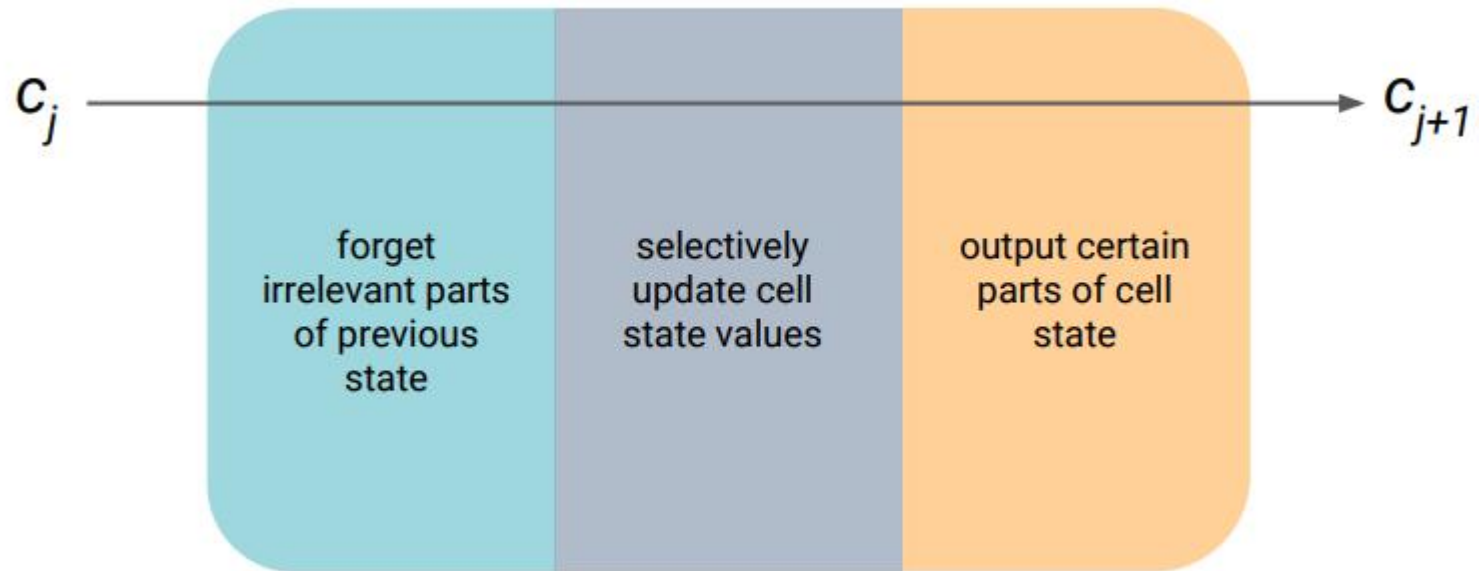
2

weights initialized to identity matrix
biases initialized to zeros


$$I_n = \begin{pmatrix} 1 & 0 & 0 & \cdots & 0 \\ 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & 1 \end{pmatrix}$$

prevents W from shrinking the gradients


More Robust Solutions: Gated Recurrent Cells



More Robust Solutions: LSTMs

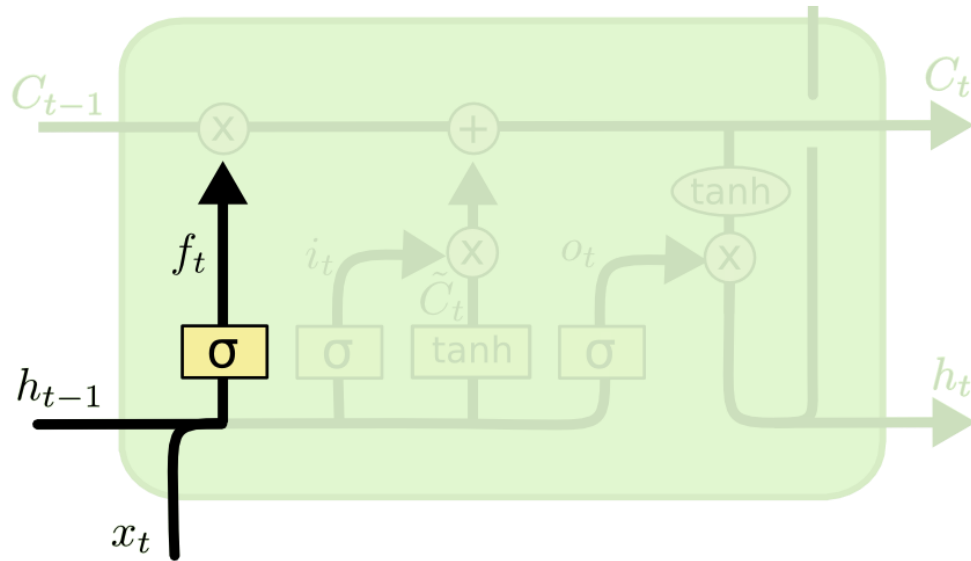


More Robust Solutions: LSTMs



More Robust Solutions: LSTMs


Forget Gate



$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

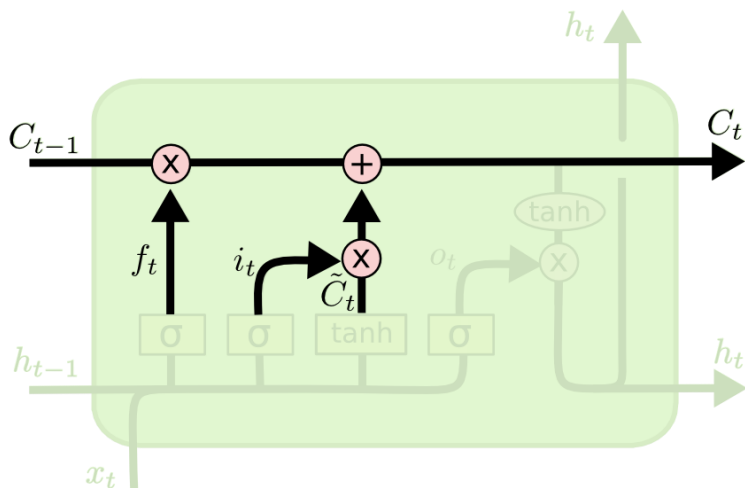
More Robust Solutions: LSTMs

Update Gate



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

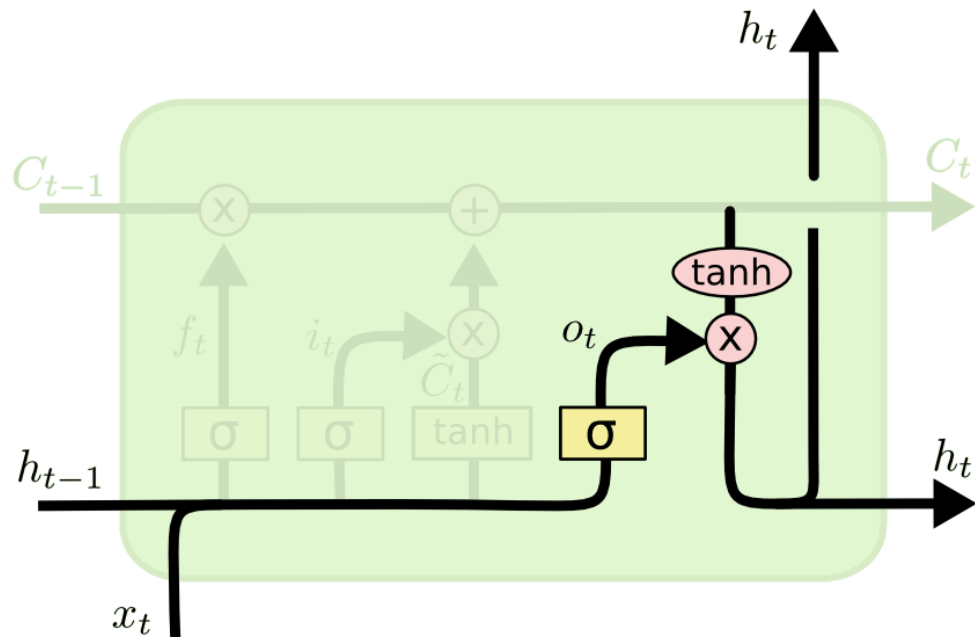
$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

More Robust Solutions: LSTMs

Output Gate



$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh (C_t)$$

Part 2 of RNNs

- **Bidirectional RNNs**
- **Attention**
- **Beam Search**
- **Coding Examples**
 - **Trigger word detection**
 - **Image Captioning(may be)**